

QBR Dashboard

Application development narration supported by requirements, user stories, application screenshots, architecture, code, Git evidence, testing, performance and business evaluation.

PREPARED

19 August 2026

NARRATION

1,992 words

EVIDENCE MODEL

Tables, figures and screenshots

TECHNICAL STATUS

20 automated tests passing

Contents

1. Project Overview
2. Requirements Analysis
3. Design and User Experience
4. Development Process
5. Testing and Quality Assurance
6. Deployment and Technical Implementation
7. Business Impact and Value
8. Critical Reflection
- References

Portfolio structure

Section	Narration focus	Principal supporting artefacts
1. Project Overview	Problem, audience, scope and value	Audience profiles and scope table
2. Requirements Analysis	Functional and non-functional needs	Requirements and prioritised user stories
3. Design and User Experience	UX rationale, accessibility and navigation	Implemented application states, user flow and mobile evidence
4. Development Process	Technology, platform and architecture appraisal	Comparison tables, architecture diagram, code evidence and Git history
5. Testing and Quality Assurance	Strategy, defects, results and code quality	Test matrix and five core evidence screenshots
6. Deployment and Technical Implementation	Distribution, infrastructure, scalability and maintenance	Deployment and scalability table
7. Business Impact and Value	Benefits, costs, alternatives and success measures	Cost-benefit, competitor and success-metric tables
8. Critical Reflection	Learning, problem solving and future practice	Evidence-linked reflective examples

1. Project Overview

Quarterly Business Reviews require information from several service areas to be converted into a concise account of performance, risk and ownership. In a spreadsheet-led process, inconsistent headings, units and status rules can create repetitive preparation and make it difficult for a reviewer to trace a warning back to its metric, owner or case. A visually convincing report can also be misleading when unlike measures are combined or when a higher value is incorrectly assumed to represent better performance.

The QBR Dashboard addresses this problem through a static browser application. It imports standardised CSV files for eight service categories, validates every row before persistence, applies direction-aware RAG logic, displays metric-separated charts and searchable records, and exports consolidated CSV and PDF reports. All processing occurs within the browser and the tested workflow made no cross-origin request. The product is intentionally a controlled prototype rather than a replacement for an enterprise business-intelligence platform.

The intended audience consists of a coordinator who prepares QBR material and a reviewer who investigates risk and accountability. These profiles are workflow-based design assumptions, not research-backed personas, because suitable participant access was unavailable within the project timeframe. Value is therefore expressed through implemented capabilities such as standardisation, traceability and local processing. Time savings, error-reduction percentages and financial returns are not claimed without a measured baseline.

Artefact 1: Constructed personas

Constructed persona 1: Sarah McLean

Attribute	Detail
Status	Fictional constructed persona; not a real Cisco employee and not interviewed
Role	Customer Delivery Architect
Experience	12 years in networking and customer delivery
Technical proficiency	Expert
Primary goals	Understand account and service health before customer business reviews; identify significant risks quickly; trace warnings to the responsible service, owner or case; produce consistent review material
Pain points	QBR information originates from several service areas; source files may use inconsistent structures; summary indicators can hide the metric that caused the warning; incorrect aggregation could lead to poor business decisions
Most relevant capabilities	Consolidated overview; RAG status; owner and case traceability; category drill-down; PDF export

Constructed persona 2: Daniel Reeves

Attribute	Detail
Status	Fictional constructed persona; not a real Cisco employee and not interviewed
Role	Technical Consulting Engineer
Experience	5 years in technical support and troubleshooting

Attribute	Detail
Technical proficiency	Advanced
Primary goals	Identify technical risks and recurring service issues; search detailed records behind QBR summaries; connect performance information with case ownership and references
Pain points	Summary data may be separated from underlying support evidence; large datasets are difficult to interrogate manually; incorrect or incomplete imported data can undermine analysis
Most relevant capabilities	Searchable detailed records; date, category and status filtering; atomic CSV validation; owner and Case ID fields; CSV export

Constructed persona 3: Priya Shah

Attribute	Detail
Status	Fictional constructed persona; not a real Cisco employee and not interviewed
Role	Network Consulting Engineer, SD-WAN
Experience	8 years in enterprise networking
Technical proficiency	Expert
Primary goals	Understand technical performance across customer environments; compare individual measures without combining incompatible units; interpret targets correctly when some metrics are higher-is-better and others are lower-is-better
Pain points	Different metrics have different units, targets and meanings; higher values cannot automatically be assumed to represent better performance; aggregating percentages, counts and durations can create misleading summaries
Most relevant capabilities	Direction-aware RAG calculation; metric-separated trend series; targets and units; category detail; boundary-safe status calculation

Scope included	Scope excluded
Local CSV import, validation and normalisation	User accounts and role-based access
Eight configured QBR categories	Central database, collaboration or multi-device sync
Charts, filters, detailed records and exports	Organisational audit history and managed backup
Local persistence with previous-state recovery	Production deployment approval and domain-owner sign-off

2. Requirements Analysis

Requirements were derived from the intended preparation and review workflow, inspection of the existing application, domain-risk analysis and defects discovered during testing. Must-have requirements focused on preserving business meaning as well as delivering visible features. This meant that invalid files had to fail atomically, status had to respect whether higher or lower was better, and unlike metrics could not be totalled simply because they were numeric.

Non-functional requirements addressed local processing, reliability, accessibility, responsiveness, maintainability and performance. Acceptance criteria were made observable so that they could be connected to automated tests or browser evidence. The requirements are implemented, but workplace terminology, target values and real metric rules still require domain-owner confirmation before operational use.

MoSCoW prioritisation separated decision-critical import, validation, status, filtering and export behaviour from supporting template and accessibility enhancements. Requirement IDs are retained in the evidence tables so implementation and test results can be traced without repeating their detail in the narration.

Artefact 2: Functional requirements and implementation status

ID	Priority	Requirement and acceptance criterion	Status	Evidence
FR-01	Must	Import a CSV into any of eight categories using replace or append mode	Complete	Upload view and browser workflow
FR-02	Must	Require Date, Category, Metric and Value headers	Complete	Parser regression tests
FR-03	Must	Reject the complete file before persistence when any material row is invalid	Complete	Figure 7 and invalid fixture
FR-04	Must	Validate real ISO dates and finite numeric values or targets	Complete	Parser tests and row-level messages
FR-05	Must	Retain explicit RAG status or derive it only from a valid target and direction	Complete	Direction and zero-target tests
FR-06	Must	Keep each metric and unit in a separate trend series	Complete	Chart tests and Figure 9
FR-07	Must	Filter inclusively by date and combine category and status filters	Complete	Filter tests and browser scenarios
FR-08	Must	Persist accepted data transactionally and recover a valid previous state	Complete	Store rollback and recovery tests
FR-09	Must	Display category summaries, charts and searchable detailed records	Complete	Overview and category-detail checks
FR-10	Must	Export complete consolidated CSV and PDF outputs	Complete	Export tests and three-engine workflow
FR-11	Should	Provide semantic data alternatives for every canvas chart	Complete	Three charts and three adjacent tables
FR-12	Should	Provide templates, clear actions and explicitly labelled demo data	Complete	Settings, empty-state and demo checks

Artefact 3: Non-functional requirements

ID	Quality attribute	Requirement	Result and limitation
NFR-01	Privacy	Process imported data without an application server	Passed in tested workflow; not an organisational security approval
NFR-02	Reliability	Preserve existing state when parsing or storage fails	Passed through atomic rejection and rollback tests
NFR-03	Accessibility	Support keyboard use, visible focus, text status and chart alternatives	Implemented; no WCAG conformance or screen-reader claim
NFR-04	Responsiveness	Remain usable at a 390-pixel viewport	Passed mobile navigation checks
NFR-05	Compatibility	Complete the core workflow in multiple browser engines	Passed Chromium 151, Firefox 153 and WebKit 26.5; WebKit is not Safari
NFR-06	Performance	Keep a 1,000-row local workflow responsive	Median import 56.4 ms, render 175.4 ms and filter 48.4 ms; baseline only
NFR-07	Maintainability	Separate parsing, storage, filters, charts, tables, upload and export	Achieved through focused IIFE modules; ordered globals remain a scaling weakness
NFR-08	Offline operation	Avoid runtime CDN dependencies	Achieved through locally bundled libraries

Artefact 4: Prioritised user stories

Every story is a design assumption because external user validation was unavailable.

Story	Priority	As a...	I need...	So that...	Acceptance evidence
US-01	Must	Coordinator	a category template	source files use the required schema	Eight template actions
US-02	Must	Coordinator	row-specific validation	source errors can be corrected before storage	Invalid fixture rejection
US-03	Must	Coordinator	replace and append modes	periodic files can be refreshed predictably	Upload mode control
US-04	Must	Reviewer	category and RAG summaries	at-risk areas can be identified quickly	KPI and status charts
US-05	Must	Reviewer	date, category and status filters	the review can focus on a relevant period or risk	Combined filter tests
US-06	Must	Reviewer	owner and case fields beside metrics	a warning can be traced to accountability	Detailed table
US-07	Must	Reviewer	separate metric trends	incompatible values are not mistaken for a meaningful total	Chart transformation tests
US-08	Must	Reviewer	complete CSV and PDF exports	evidence can be shared without re-keying	Export regression tests
US-09	Should	Keyboard user	visible focus and controllable mobile navigation	the core interface is operable without a pointer	Keyboard browser checks
US-10	Should	Assistive-technology user	structured chart-data tables	values are not available only as canvas pixels	Figure 9 and semantic tables

3. Design and User Experience

The design uses a summary-to-evidence hierarchy. The overview presents independent category status, category detail separates metric trends and RAG distribution, and the record table exposes dates, targets, owners and case references. This supports recognition and progressive disclosure instead of showing the full dataset at once. A persistent desktop sidebar becomes an overlay drawer at narrow widths, retaining the same category and workspace structure rather than introducing a different mobile information architecture.


Error prevention received greater priority than decorative complexity. Invalid imports report CSV row numbers and reasons, preserve existing data and reject the entire file before any partial chart is shown. Templates and the schema reference reduce memory demand, while explicitly labelled demonstration data avoids confusion with live information. These decisions apply visibility, consistency, prevention and recovery principles associated with heuristic evaluation (Nielsen, 1994).

The colour palette uses deep navy for navigation, blue for actions and labelled red, amber and green states. RAG meaning is never communicated by colour alone. Typography uses the system font stack to avoid an external font request and to retain familiar platform rendering. Accessibility decisions target keyboard operation, focus visibility, labels, status messages and reflow in WCAG 2.2 (W3C, 2024). Since Chart.js canvas content is not inherently exposed to screen readers, each named canvas has an adjacent captioned data table (Chart.js, 2025).


Representative testing could have identified terminology, recovery or interpretation problems that expert inspection cannot reliably detect. Participant access limitations prevented this work, so no user quotations, ratings, task success or feedback-led iteration are claimed. The captured application states and implemented changes are presented as design and test evidence rather than user-validated research.

Low-fidelity design wireframes

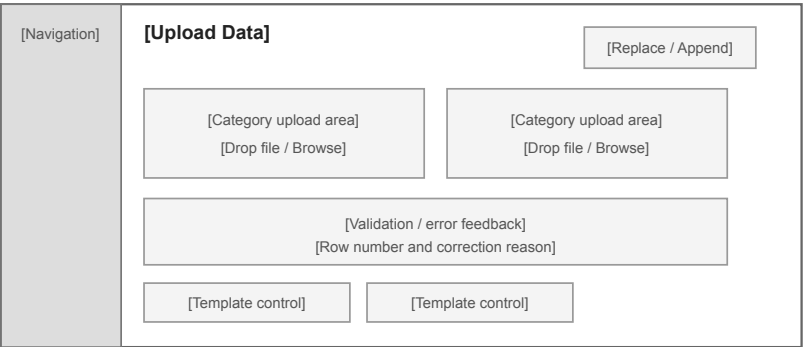
1. Overview Dashboard




2. Category Detail



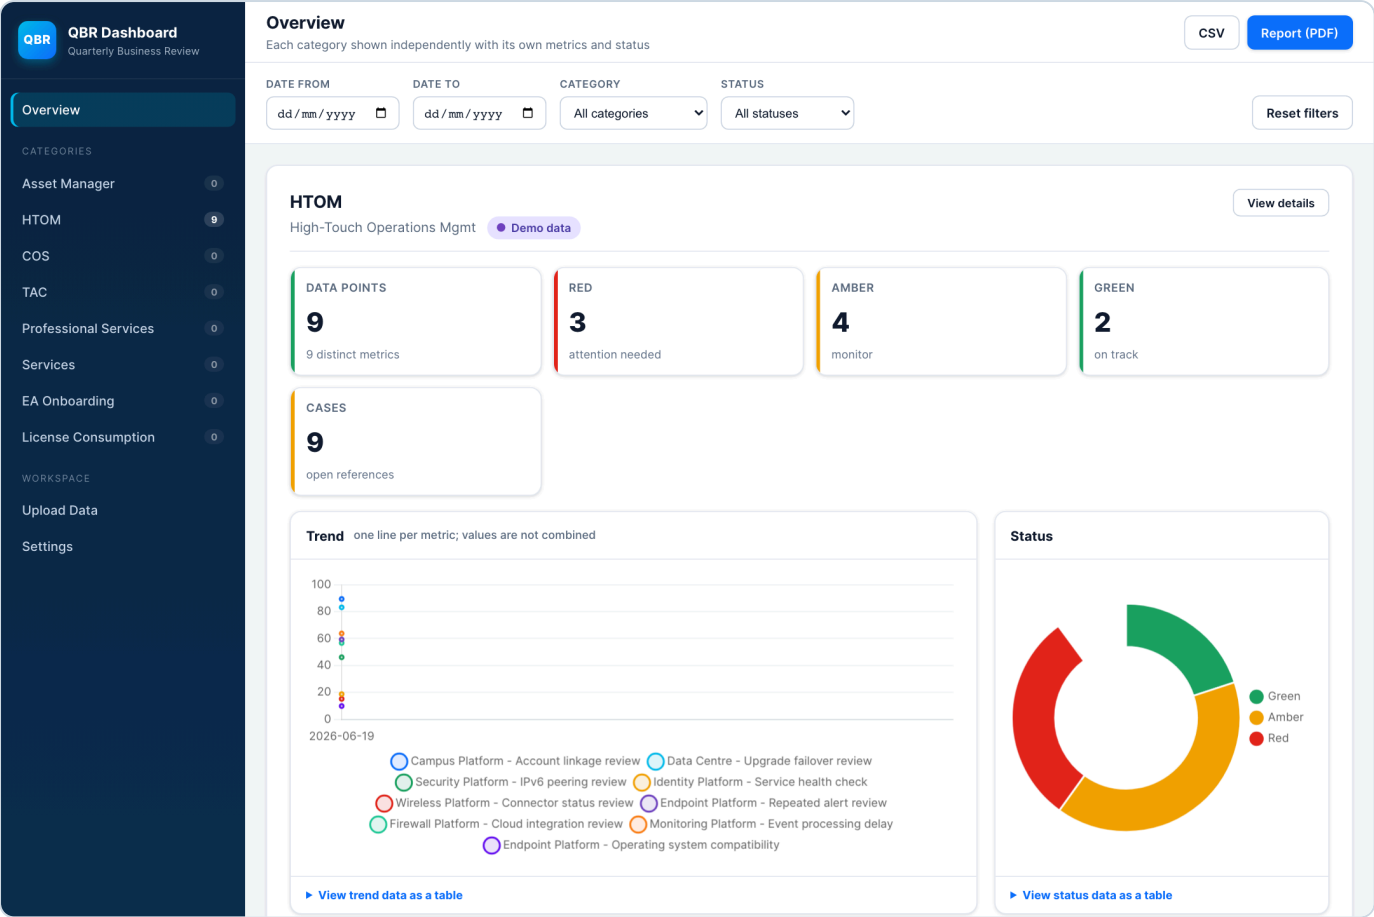
3. Upload Data



4. Mobile Category View



Artefact: Low-fidelity design wireframes documenting the dashboard information architecture. These retrospective drawings document the structure represented by the implemented application; they are not claimed as pre-development evidence.



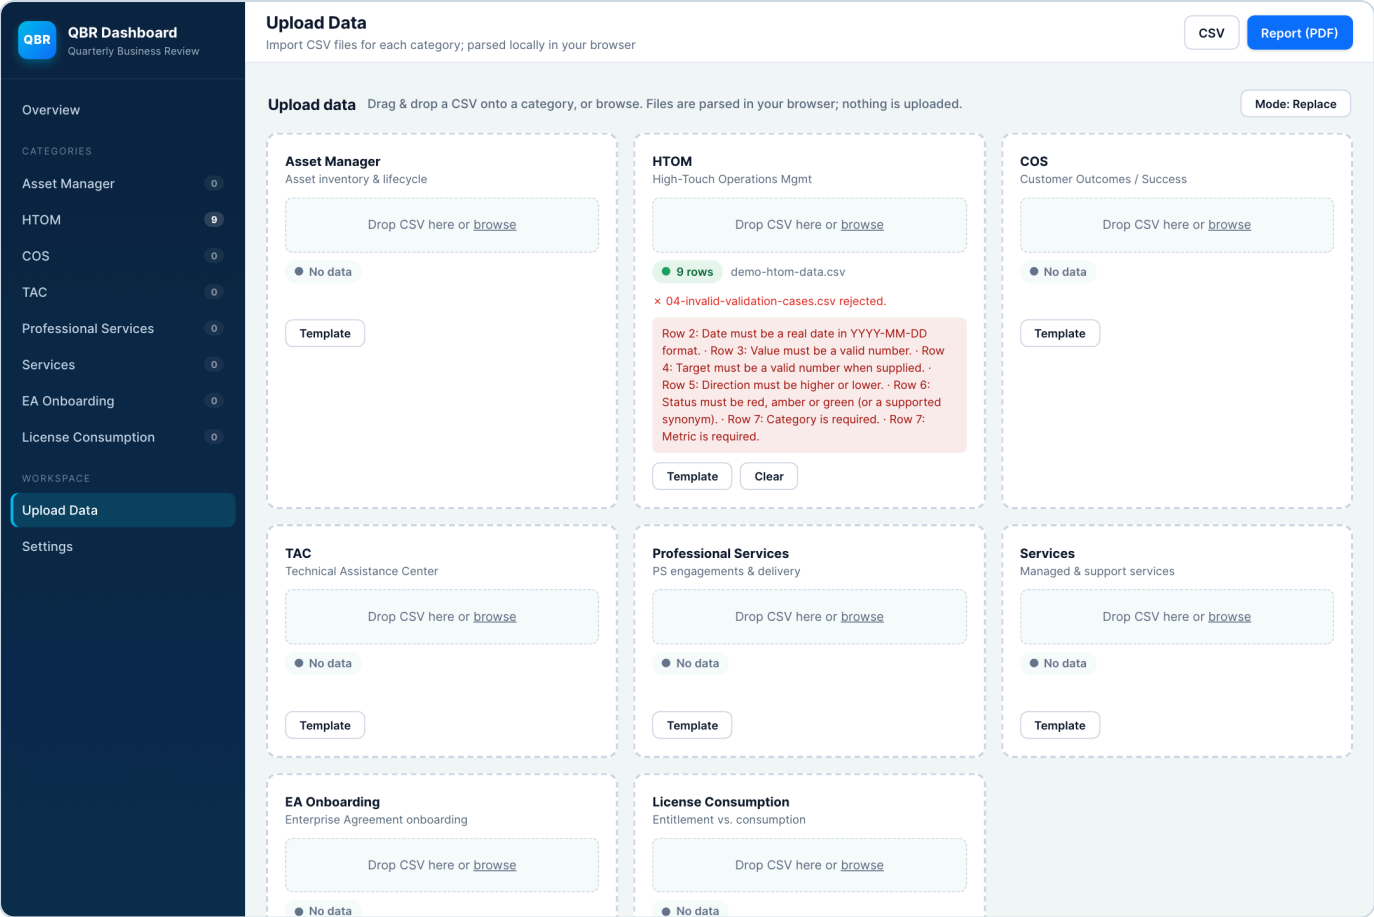


Figure 1b. Implemented upload validation showing the invalid fixture rejected while the existing nine HTOM rows are preserved.

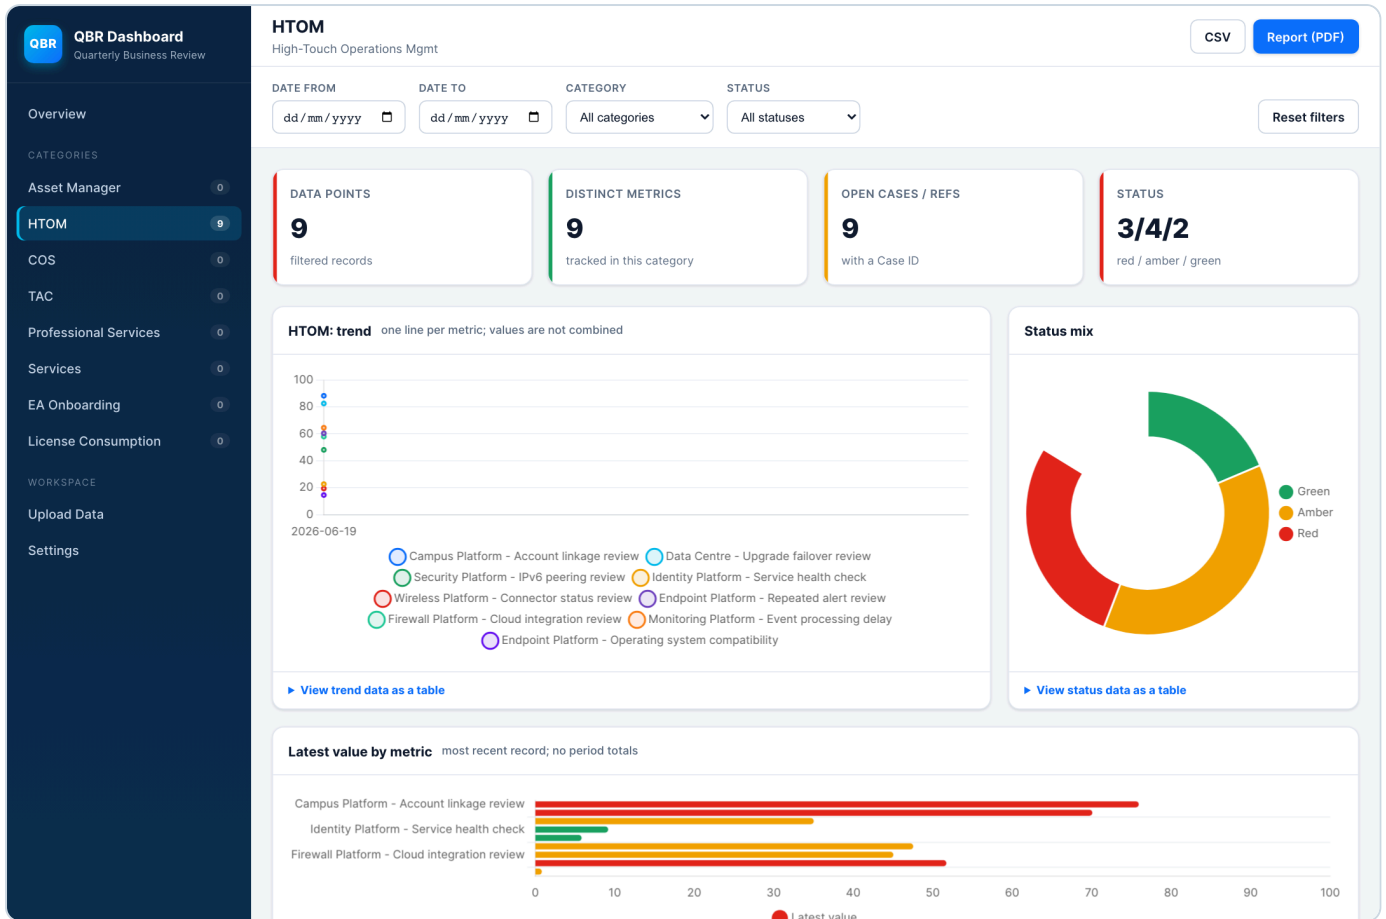


Figure 1c. Implemented category detail showing summary measures, metric trends, status distribution and latest values.

HTOM

High-Touch
Operations Mgmt

CSV

Report (PDF)

DATE FROM

DATE TO

dd/mm/yyyy

dd/mm/yyyy

CATEGORY

STATUS

All categories

All statuses

Reset filters

DATA POINTS

9

filtered records

DISTINCT METRICS

9

tracked in this category

OPEN CASES / REFS

9

with a Case ID

STATUS

3/4/2

red / amber / green

HTOM:

one line per metric; values are not

Figure 1d. Implemented category detail at a 390-pixel viewport, with controls and summary measures reflowed for mobile use.

Core preparation and review workflow




Figure 2. Core preparation and review flow, including correction and re-import after atomic rejection.




Figure 3. Responsive navigation at 390 by 844 pixels. Keyboard opening, focus entry, Escape closure and focus return were also checked across three engines.

4. Development Process

A static web platform was selected because the workflow is file-based, requires no shared live database and must run on common desktop operating systems. HTML supplies semantic structure, CSS handles responsive presentation and JavaScript performs validation, transformation and interaction. The same source can be served on macOS, Windows or Linux, while locally bundled libraries remove a runtime internet dependency.

Vanilla JavaScript was proportionate to the number of views and absence of a remote API. React or Vue could offer component boundaries, build tooling and stronger large-team conventions, but would add dependency and migration overhead without resolving the main domain risks. The IIFE modules separate concerns clearly, although reliance on global objects and ordered script loading would become a maintainability constraint if the product gained concurrent developers, remote services or frequent releases.

Specialist libraries were used where rebuilding mature behaviour would increase risk. PapaParse handles quoted CSV fields and tolerant parsing; Chart.js provides responsive charts; DataTables supplies search, sort and pagination; and jsPDF with AutoTable creates client-side reports. This shortens development but transfers responsibility for dependency review, lifecycle management and accessibility adaptation. DataTable and Chart instances are therefore destroyed before rerendering, and canvas charts receive structured alternatives.

The client-side platform was compared with native, hybrid, hosted and business-intelligence alternatives. A native desktop product would improve operating-system file integration but add packaging and update channels. A mobile application would not suit dense report preparation. A hybrid wrapper would reuse web code but add build complexity. A hosted service or BI platform would provide authentication, collaboration and central governance, but would require approved infrastructure, data-transfer controls and potentially licensing or vendor lock-in. The selected approach is therefore appropriate for a bounded prototype, not universally superior.

Development became increasingly test-led. Initial logic assumed higher values were better, so incident-style measures and zero targets could receive false green status. The schema gained an explicit Direction field and safe boundary logic. Mixed metrics were removed from totals and keyed into separate series. Storage updates became transactional with previous-state recovery. The PDF row limit was removed, spreadsheet-formula-like CSV values were neutralised, and demo data became explicit opt-in content. These changes demonstrate progression from visible functionality to semantic correctness and recoverable failure.

QBR Dashboard: client-side architecture




Figure 4. Client-side module architecture. The application orchestrator coordinates focused modules, while accepted state is persisted to a primary and previous-state browser key.

Artefact 5: Platform and framework appraisal

Option	Strengths for this context	Weaknesses for this context	Decision
Static web application	Portable, inspectable, offline-capable and no installation package	Origin-specific local storage and no central governance	Selected for the bounded prototype
React or Vue web application	Components, ecosystem and scalable build tooling	Build and dependency overhead for a small local application	Deferred until complexity justifies migration
Native desktop	Strong file and OS integration	Separate packaging, signing and update distribution	Rejected for current scope
Native mobile or hybrid	Touch support and possible app-store distribution	Dense QBR tables are desktop-oriented; wrapper complexity	Rejected for primary workflow
Hosted web service	Authentication, backup, collaboration and auditability	Infrastructure, security, governance and hosting requirements	Appropriate if shared operation becomes mandatory
BI platform	Managed analytics and organisational integration	Licensing, vendor lock-in and reduced control over unique validation rules	Operational alternative, not selected for the development portfolio

Artefact 6: Representative code decision

The status function refuses to guess when target or direction is absent and handles zero targets explicitly.


```
function deriveStatus(value, target, direction) {
  if (value == null || target == null || !direction) return '';
  if (direction === 'higher') {
    if (target === 0) return value >= 0 ? 'green' : 'red';
    if (value >= target) return 'green';
    return value >= target * 0.8 ? 'amber' : 'red';
  }
  if (direction === 'lower') {
    if (target === 0) return value <= 0 ? 'green' : 'red';
    if (value <= target) return 'green';
    return value <= target * 1.25 ? 'amber' : 'red';
  }
  return '';
}
```

Artefact 7: Version-control and CI evidence

Evidence	Verified fact	Critical interpretation
59d679d	Initial project import	Establishes the repository baseline
0e316b5	Dashboard scaffold commit	Shows a separate scaffold stage
cbb41b1	Reconstructed-history branch merged into <code>main</code>	History exists but is less granular than an ideal requirement-linked sequence
Final release commit	df32727b765e5bea64f972e0ca83244672ea502c on <code>main</code>	Records the application, portfolio and automated release publication state
GitHub Actions workflow	<code>QBR Dashboard Tests</code> , run 3, completed successfully on 19 August 2026	The remote <code>Run regression tests</code> step passed for the final release commit; the local suite reported 20/20
Release evidence	Annotated tag <code>v1.1.0</code> , <code>CHANGELOG.md</code> and GitHub Release published with Word and PDF assets	Improves traceability, but the historical commit sequence remains less granular than an ideal requirement-linked history

Jack-Mitchell27 / QBR-Dashboard-Web-DevPublic

NotificationsFork 0Star 0

<> CodeIssuesPull requestsActionsProjectsSecurity and qualityInsights

← QBR Dashboard Tests

✓ Add automated release publication #3Sign in to view logs

Summary

All jobs

test

Run detailsUsageWorkflow file

Triggered via push 2 minutes ago

Status

Total duration

Artifacts

Jack-Mitchell27 pushed -> df32727main

Success

11s

-

test.yml

on: push

test7s

Figure 5. GitHub Actions run 3 completed successfully for commit df32727 on main. The test job passed remotely before the v1.1.0 release was published.

Testing combined pure-logic regression tests with browser workflows, negative fixtures, accessibility scanning and performance measurements. The Node.js test runner was suitable because it is stable, dependency-light and returns a failing process status when assertions fail (Node.js contributors, n.d.). Twenty tests cover parser boundaries, metric transformations, filters, transactional storage, backup repair and export completeness. Browser checks then verify rendering, downloads, focus and library integration that the pure harness cannot exercise.

The most valuable tests challenged meaning rather than appearance. Invalid files were rejected atomically; lower-is-better and zero-target status branches were verified; unlike units remained separate; failed persistence preserved previous state; and a 45-row export helper proved that the historical PDF limit had been removed. CSV fields beginning with spreadsheet formula markers are prefixed as a risk reduction, while OWASP guidance is acknowledged that no mitigation is universal for every spreadsheet consumer (OWASP Foundation, n.d.).

The same 48-row workflow passed in Chromium 151, Firefox 153 and WebKit 26.5 with no captured page error, console error or external request. Twelve axe-core scans initially found supporting-text and transient toast contrast failures. Both were corrected, after which final scans reported zero automatically detected violations. This does not establish WCAG conformance. Five Chromium runs with 1,000 rows established local medians of 56.4 ms for import, 175.4 ms for detail rendering and 48.4 ms for filtering.

Artefact 8: Test strategy and results

Level	Scope	Quantity	Result	Limitation
Parser regression	Headers, dates, numbers, enums, direction and zero targets	7 tests	Pass	Pure harness, not browser file control
Chart regression	Metric separation, duplicates, latest values and RAG roll-up	4 tests	Pass	Canvas rendering checked separately
Filter regression	Date boundaries, combined filters and reset	3 tests	Pass	Uses representative fixtures
Store regression	Persistence, rollback, recovery and clear-all	4 tests	Pass	In-memory storage harness
Export regression	CSV protection and complete PDF body mapping	2 tests	Pass	File download checked in browsers
Browser integration	End-to-end views and actions	15 scenarios	Pass	Controlled workflow, not organisational acceptance
Accessibility automation	Four views across three engines	12 scans	Zero final detected violations	Manual screen-reader testing unavailable
Performance	1,000-row import, render and filter	5 runs	Responsive local baseline	No agreed production threshold

QBR Dashboard regression evidence

Automated unit and integration checks completed successfully.

✓ ALL TESTS PASS

TESTS

20

PASSED

20

FAILED

0

Verified test groups

✓ Charts and metric separation	4 / 4
✓ CSV and PDF export	2 / 2
✓ Combined filtering	3 / 3
✓ Parsing and status logic	7 / 7
✓ Persistence and recovery	4 / 4

\$ npm test

Node.js 22.14.0 | 20 tests | 0 failures

Figure 6. Automated regression evidence: 20 tests passed with zero failures or skips.

Services

Managed & support services

Drop CSV here or [browse](#)

● No data

× 04-invalid-validation-cases.csv rejected.

Row 2: Date must be a real date in YYYY-MM-DD format. · Row 3: Value must be a valid number. · Row 4: Target must be a valid number when supplied. · Row 5: Direction must be higher or lower. · Row 6: Status must be red, amber or green (or a supported synonym). · Row 7: Category is required. · Row 7: Metric is required.

Template

Figure 7. Negative testing: invalid date, value, target and direction data was rejected before persistence.

Show 10 ▾

Search...


DATE ▴ ▾	CATEGORY ▴ ▾	METRIC ▴ ▾	VALUE ▴ ▾	TARGET ▴ ▾	DIRECTION ▴ ▾	STATUS ▴ ▾
2026-02-28	Incidents	Open P1 Cases	2	0	lower	● Red
2026-01-31	Incidents	Open P1 Cases	0	0	lower	● Green

Showing 1–2 of 2

Previous 1 Next

Figure 8. Boundary testing: zero meets a zero lower target, while a positive value against it is red.

Services: trend one line per metric; values are not combined



▼ **View trend data as a table**

Services metric values by date. Metrics may use different units and are not combined.

Metric	2026-01-31	2026-02-28	2026-03-31	2026-04-30
Incidents - Open P1 Cases (count)	0	2	No record	No record
Support - Resolution Time (hours)	8	12	13	No record
Operations - Backlog (count)	15	20	24	26

Figure 9. Unlike metrics remain separate and the chart has an expandable semantic data alternative.

Show10

March

DATE	CATEGORY	METRIC	VALUE	TARGET	DIRECTION	STATUS
2026-03-31	Service Quality	Availability	79	99	higher	● Red
2026-03-31	Customer Experience	CSAT	65	85	higher	● Red
2026-03-31	Incidents	Open P1 Cases	3	2	lower	● Red
2026-03-31	Support	Resolution Time	11	8	lower	● Red
2026-03-31	Assets	Assets Patched	70	95	higher	● Red
2026-03-31	Licensing	Licences Used	60	80	higher	● Red
2026-03-31	Operations	Backlog	28	20	lower	● Red
2026-03-31	Performance	SLA Met	75	98	higher	● Red

Showing 1–8 of 8 (filtered from 48 total entries)

Previous

1

Next

Figure 10. Searching the 48-row fixture for March returned the expected eight detailed records.

QBR Dashboard: Expanded QA Evidence

Cross-browser workflow, accessibility scanning and 1,000-row performance

18 August 2026
macOS, Apple silicon

Browser workflow

Chromium	151.0.7922.34	PASS
Firefox	153.0	PASS
WebKit	26.5	PASS

Each engine verified clean start, 48-row import, 3 charts, 3 data alternatives, 16-row red filter, CSV/PDF export, mobile keyboard navigation and skip-link behaviour.

Chromium performance, 5 runs

Import 1,000 rows	56.4 ms	54.6 to 57.9 ms
Render detail	175.4 ms	173.4 to 176.6 ms
Apply status filter	48.4 ms	47.1 to 55.3 ms
Page load	61.1 ms	59.8 to 74.1 ms

Median shown with observed minimum and maximum. Local loopback test on one development machine, not a production service level.

20/20

automated regression tests passed

12

final axe-core view and engine scans

0

automatically detected violations after fixes

0

page errors, console errors or external requests

Accessibility defects found and corrected

ID	INITIAL FINDING	CORRECTION	RETEST
A11Y-01	Six text nodes measured 3.06:1 on white	Faint text changed from #8595a8 to #64748b	PASS
A11Y-02	Toast animation produced transient 3.75:1 contrast	Opacity removed from the entrance animation	PASS

Evidence boundary: WebKit is a compatibility proxy, not genuine Safari. Automated scans do not prove WCAG conformance. Manual contrast, screen-reader, representative-user, target-device and organisational security checks remain.

Figure 11. Three-engine workflow results, corrected contrast defects and five-run performance medians.

6. Deployment and Technical Implementation

The application is distributed as static files with platform launchers and no production build step. A local HTTP origin is preferred because `localStorage` is origin-specific and direct file behaviour varies between browsers (MDN contributors, 2026). Vendor libraries are bundled, so runtime operation does not require a CDN. The recipient can use the supplied macOS, Windows or Linux launcher, while the Quick Start document explains import, backup and troubleshooting.

This approach scales operationally only to the storage, processing power and browser profile of one device. Increasing row volume can be measured, but increasing users creates isolated copies rather than a shared service. A hosted architecture would be required for authentication, central backup, audit history and concurrent use. Maintenance currently involves dependency review, regression execution, browser checks and clear release notes. The final `main` commit passed the remote GitHub Actions regression workflow before the annotated `v1.1.0` tag and GitHub Release were published.

Artefact 9: Deployment, scalability and maintenance

Concern	Current implementation	Scaling or maintenance implication
Distribution	Static folder and three platform launchers	Simple handover; no automatic updater
Hosting	Loopback HTTP server	No external hosting service in tested mode
Data storage	Browser <code>localStorage</code> plus previous-state backup	Device and origin specific; no shared recovery
Runtime dependencies	Libraries bundled in <code>vendor</code>	Offline-capable, but versions require periodic review
More data	Browser-side parse, render and storage	Re-run thresholds on target device and maximum expected volume
More users	Independent browser profiles	Requires hosted authentication and central data services
Release quality	Final <code>main</code> commit <code>df32727</code> ; local 20/20 tests; successful remote workflow run 3; annotated <code>v1.1.0</code> tag; changelog and GitHub Release	Verified release evidence, while production approval and stakeholder acceptance remain unavailable

7. Business Impact and Value

The dashboard's immediate value is structural rather than a proven financial return. It standardises the input schema, rejects materially invalid files, preserves metric direction and unit, connects summary status to detailed ownership, and produces consistent exports. Local processing also reduces application-driven data transfer during the tested workflow. These capabilities can support faster preparation and fewer interpretation errors, but those outcomes were not measured and are therefore framed as hypotheses.

The main alternative is a spreadsheet-based process, which is flexible and familiar but can duplicate validation, formatting and chart work. A commercial BI platform provides stronger collaboration, governance and support but can add licensing, configuration and vendor dependence. A custom hosted service would improve multi-user operation but require security, authentication and infrastructure. The QBR Dashboard occupies a low-infrastructure prototype niche: more controlled than manual spreadsheet assembly but less governed than an enterprise platform.

A cost-benefit assessment cannot responsibly assign cash values without current preparation time, salary rates, error costs, usage frequency and maintenance effort. The table therefore distinguishes observed implementation facts from measurements needed in a future organisational evaluation. This is stronger than presenting unsupported savings.

Growth potential depends on governance needs. Confirmed demand for shared access could justify authentication, central storage and managed deployment, while approved data connectors could reduce manual file preparation. These investments should follow evidence of adoption rather than being added speculatively to the prototype.

Artefact 10: Cost-benefit framework without invented values

Factor	Current cost or benefit evidence	Monetary estimate	Data required for valuation
Development	Existing project work and maintenance time	Not measured	Actual hours and agreed labour rate
Local hosting	No separate hosting service in tested mode	No external hosting charge evidenced	Device, support and operational overhead
Libraries	Bundled project dependencies	No paid licence cost documented	Formal licence and procurement review
Maintenance	Regression, browser and dependency checks required	Not measured	Release frequency and support ownership
Preparation efficiency	Consolidated import, filters and exports implemented	Benefit not measured	Before/after task timing
Error reduction	Atomic validation and semantic tests implemented	Benefit not measured	Baseline and post-adoption error counts
Traceability	Owner, case, target and metric retained together	Qualitative benefit evidenced	Operational audit requirements
Data transfer	No external request captured in tested workflow	Risk reduction, not cash saving	Organisational security valuation

Artefact 11: Alternative solutions

Alternative	Advantage	Disadvantage compared with this prototype
Manual spreadsheet workbook	Familiar and highly flexible	Repeated validation, aggregation and presentation remain manual
Commercial BI platform	Collaboration, governance and vendor support	Licensing, setup and less control over bespoke CSV rules
Custom hosted application	Central data, authentication and audit history	Infrastructure, security and ongoing service ownership
QBR Dashboard	Local, portable, transparent and tailored validation	No central collaboration, governance or validated user outcomes

Artefact 12: Success measures and current evidence

Measure	Method	Current result	Interpretation
Functional regression	Automated test pass rate	20/20 passed	Strong repeatable logic evidence
Invalid-import safety	Persisted rows after invalid fixture	0 rows	Atomic rejection achieved
Browser workflow	Engines completing standard scenario	3/3	Chromium, Firefox and WebKit proxy passed
Automated accessibility	Final detected violations	0 across 12 scans	Positive automated result, not conformance
1,000-row import	Five-run median	56.4 ms	Responsive on recorded development machine
1,000-row detail render	Five-run median	175.4 ms	Responsive on recorded development machine
User task success	Representative observation	Not measured	Participant access unavailable
Preparation time or savings	Before/after workplace baseline	Not measured	No efficiency or ROI claim permitted

8. Critical Reflection

The most significant learning was that visible completeness is not the same as decision quality. Early development concentrated on upload, charts, tables and exports, but testing exposed that the application could display a polished yet incorrect conclusion. The higher-is-better assumption suited availability but reversed measures such as incident count or resolution time. A positive value against a zero lower target could even become green. Introducing Direction, refusing to derive status without a valid rule and testing zero-target boundaries shifted the focus from presentation to semantic correctness.

A similar lesson came from aggregation. Combining numeric values looked convenient, yet percentages, hours and counts do not form a meaningful total. Replacing the total with a data-point count and keying trend series by metric and unit demonstrated that mathematically valid operations can still be analytically wrong. Future projects should define the unit, direction, target owner and worked examples for each measure before visualisation begins.

Failure evidence improved the process. Forced storage and corrupt-state cases led to rollback and backup repair, while export tests exposed the historical row limit and formula risk. Accessibility scans found two contrast defects, which were corrected and rescanned. An apparent WebKit focus failure proved to be a pointer-based test-method issue; repeating the interaction with keyboard focus and Enter verified the intended behaviour. The test method therefore required scrutiny alongside the code.

The modular structure and dependency-free Node tests made domain rules easier to isolate, while multi-engine checks covered integration behaviour. However, browser-heavy branches remain less covered by the pure harness, the historical commit sequence is less granular than an ideal requirement-linked history, and expert evaluation cannot replace user observation. Representative testing would have improved confidence in terminology, recovery and interpretation, but access was unavailable. Recording this limitation is more professionally credible than inventing feedback.

If repeated, semantic criteria and fixtures would precede interface work, commits would be smaller and linked to requirement IDs, and deployment constraints would be agreed earlier. Quality is now understood as correct meaning, recoverable failure, accessible operation, traceable evidence and explicit limits. This makes residual risk visible to technical and business stakeholders.

References

Chart.js (2025) *Accessibility*. Available at: <https://www.chartjs.org/docs/latest/general/accessibility.html> (Accessed: 18 August 2026).

MDN contributors (2026) *Window: localStorage property*. MDN Web Docs. Available at: <https://developer.mozilla.org/en-US/docs/Web/API/Window/localStorage> (Accessed: 18 August 2026).

Nielsen, J. (1994) 'Enhancing the explanatory power of usability heuristics', in *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*. New York: ACM, pp. 152-158. doi: 10.1145/191666.191729.

Node.js contributors (n.d.) *Test runner*. Node.js documentation. Available at: <https://nodejs.org/api/test.html> (Accessed: 18 August 2026).

OWASP Foundation (n.d.) *CSV Injection*. Available at: https://owasp.org/www-community/attacks/CSV_Injection (Accessed: 18 August 2026).

World Wide Web Consortium (W3C) (2024) *Web Content Accessibility Guidelines (WCAG) 2.2*. W3C Recommendation, 12 December. Available at: <https://www.w3.org/TR/WCAG22/> (Accessed: 18 August 2026).